


Qwen3-4B / Layer 27

(a) Fixed prompt length $m = 4$ · $n = 95$



Mean $\pm$ std across inputs · raw, unnormalized attention-output differences

(b) Single-token steering · $n = 400$

